



Operator's Manual

For your hardened OpenClaw agent. Read what you need, skip what you don't.

*The most loved manuals don't explain a product.
They teach you how to operate in a new world.*

WELCOME

The three things that matter most

This is the manual for **Juno**, your hardened OpenClaw agent.

You bought a setup, not an AI. The AI is the easy part. What's hard is installing it well — sandboxed, opinionated, version-controlled, with sensible defaults and a manual you can actually read. That's what's in this document.

Read what you need, skip what you don't. The Quick Reference Card on the next spread is for daily use. The deeper sections are for when something surprises you.

If you remember nothing else from this manual, remember these three things:

- ◆ **Less, but tested.**

Juno does less than other agents on purpose. The restraint is the product.

- ◆ **The friction is the feature.**

Juno will ask you before doing things that matter. That's not annoyance; that's safety with stakes attached.

- ◆ **You can always say "stop everything."**

Five seconds. Halts every running task. No questions asked.

The rest of this manual is detail. Start with the Quick Reference Card. Come back here when you need it.

QUICK REFERENCE CARD

Print this. Stick it where you can see it.

EMERGENCY

If this happens	Text Juno
Something feels off	<code>stop everything</code>
You want it back	<code>resume normal operations</code>
You want to know what's running	<code>show all running tasks</code>
You want to override one rule once	<code>I approve this specific instance, knowing it bypasses [rule]</code>

DAILY

What you want	What to send
Yesterday's spend	<code>Arrives every morning at 8am Rome time</code>
A specific email summary	<code>Summarize emails from [sender] this week, two sentences each</code>
Update what Juno knows about you	<code>Update USER.md to say [X]</code>
Pause one scheduled task	<code>Pause the [task name] job</code>

WHEN THINGS BREAK

1. **No response?** Send "hello?" If still nothing, check the Hetzner status page.
2. **Spending spiked?** Text why was yesterday so expensive?
3. **Cryptic error?** Screenshot. Email us.
4. **Juno acted unexpectedly?** Read the relevant AGENTS.md section (text show me AGENTS.md §X). The rule is there.

HARD LIMITS AT A GLANCE

Limit	Value
Daily spend, soft warning	€1 (alerted in morning report)
Daily spend, hard refuse	€4 (configurable)
Outbound messages	3/hour, always asks first
Reading/research	30 actions/hour
Auto-pause on silence	14 days (configurable: 7/14/30/off)

CONTACT

Email	junotheclawagent@gmail.com
During 30-day warranty	Anything. 24-hour response.
On retainer	Anything. 24-hour response.
Off retainer	Emergencies only. 3-5 business days.



PART I

Meet Juno

Time to introduce the actual thing you bought.

The four sticky notes on Juno's desk

Imagine you've hired a new assistant who'll work with you for years. On their first day, you sit them down and explain how you want them to talk, the rules of the job, who's who in your life, and what they should know about you.

You write all four down so they don't forget. You stick the notes on their desk.

That's exactly what we did for Juno. Four short documents, sitting in a folder on your server, that load every time the agent wakes up:

SOUL.md — How Juno talks

The voice. Direct, no emoji unless you use one first, no "Great question!" openers, no fake enthusiasm. Writes like a senior colleague — confident enough to push back, warm enough to be liked. Says less, not more. Doesn't fill silence.

This file is what makes Juno feel different from ChatGPT. Don't edit it in the first month. Live with the voice we set; tune only if there's something specific you'd change.

AGENTS.md — What Juno will do and won't

The rulebook. Longest of the four files, and the most important. Contains the **Ask First** list, the **Never** list, web browsing rules, privacy rules, spending limits, kill switch behavior, and weekly reflection schedules.

Don't edit without thinking. Each rule is there because something specific could go wrong without it.

IDENTITY.md — Who Juno knows

A short document establishing who you are (the principal), who Juno is (the agent), and the boundaries between you. Includes the hard limits on impersonation — Juno will never represent itself as you in romantic, legal, medical, financial, or family correspondence. Ever.

USER.md — What Juno knows about you

The things you'd tell a new assistant on day one: your name, where you live, your job, your projects, your preferences, your email signature, the languages you speak.

A useful habit: when Juno gets something wrong about you, the answer is usually *"edit USER.md."* Tell Juno what to change. It edits the file and confirms.

Where Juno lives

A small computer in an EU country of your choosing — by default, Helsinki, Finland (Hetzner CX23 server, 4 GB RAM, 2 cores, 40 GB disk).

Costs **€4.87/month, billed by Hetzner in Germany. Includes German VAT; deductible if you're VAT-registered.** Plus €0.76/month for automatic daily backups, which we enabled.

Why EU

Your data physically stays in the European Union, under GDPR — the strictest data protection law in the world. Not American jurisdiction. Not Chinese jurisdiction. EU.

That matters because if anything ever goes wrong — a breach, a dispute, a question about your data — you have the strongest possible legal framework standing behind you. We didn't put your data in Helsinki by accident.

What you can do with it

Shut it off anytime. Log into Hetzner dashboard, click "Stop." Juno goes offline. You stop paying for API usage. Click "Start" — back online.

Wipe it entirely. Same dashboard, "Delete server." Gone in 60 seconds. Your data, your agent, everything. We can't stop you. We also can't undo it.

Nobody else can reach it. Firewall allows only two things in: SSH (for you and us to log in) and the agent's internal gateway (only listens on the server itself, never the internet).

What gets remembered, and what gets thrown away

A common worry: *"Is this thing recording my entire life and saving it forever?"*

No. Here's the honest accounting:

What	Where it lives	How long
Conversations with Juno	Your server, plain text	Auto-reset every 7 days (configurable)
Memory file (USER.md)	Your server	Until you delete a line yourself
Email contents (if Gmail connected)	Nowhere	Read in the moment, released from memory
Logs	Server, system folder	Auto-deleted after 30 days
Daily spend reports	Your server	As long as you want them
Trajectory files (Juno's thinking)	Your server	Deletable in one command
Phone calls, video, audio, location	—	Never recorded

Juno only sees what comes through Telegram and what's in folders you specifically granted it.

Version control: your safety net

Every change Juno makes to its own cognitive files gets recorded in a version history — like Track Changes in a Word document, but automatic. Every night at 3am, the server takes a snapshot of all four files.

You can never permanently lose a version of Juno.

To roll back a change:

EXAMPLE CONVERSATION

You: Show me yesterday's version of USER.md.

Juno: [shows it]

You: Restore that version.

Juno: [restores; commits the change to git history]

That's the whole flow. You don't need to know what git is. Juno handles it.

TRY IT NOW

- Show me what you know about me from USER.md.

It'll summarize what's in your memory file — name, location, projects, preferences. That's what Juno has on its desk. Anything wrong or missing? Tell Juno to fix it.



PART II

Day-to-day

The single biggest thing that separates people who get value out of Juno from people who don't: **how they phrase requests.**

Three habits. Once they're habits, you don't think about them.

Habit 1: Be specific about scope

Vague	Specific
<i>"summarize my emails"</i>	<i>"summarize emails from my team this week — two sentences each, group by sender"</i>
<i>"help me with the article"</i>	<i>"draft three opening paragraphs for the al-Khwarizmi chapter, narrative voice, around 200 words each"</i>
<i>"what's on my mind today"</i>	<i>"summarize my calendar today plus three open Sentra threads I should follow up on"</i>

The vague version makes Juno guess. The specific version makes it work.

If you get back something you didn't want, the right move isn't to give up — it's to refine in the moment: *"too long, give me three bullets each"* or *"only emails about the book project."*

Habit 2: Tell it the format you want

Juno will pick a format if you don't specify. Often the wrong one. Just ask:

- *"Give me a table — three columns: sender, subject, what they want."*
- *"One paragraph. No bullet points."*
- *"A short list, no more than five items."*

This is the single fastest way to make Juno feel "right" for what you're doing.

Habit 3: When it's wrong, say so

Juno won't take offense. If it gave you a five-paragraph answer and you wanted two sentences, say *"two sentences."* If it used emoji and you didn't want any, say so.

If Juno gets the same thing wrong twice, say: *"Remember in USER.md that I prefer X."* It edits the file. Next session, no more correction needed.

The morning spend report

Every morning at 8am Rome time, a message arrives in Telegram:

Yesterday's spend · May 14, 2026

AI model used: Claude Haiku 4.5

Total cost: €4.07 (\$4.46)

Status:  Above €1 soft cap – heavy product-building day

Words you sent: ~650 (input: 860 tokens)

Words Juno wrote back: ~49,000 (output: 65,344)

Stuff Juno remembered (cheap): ~7.3M words (cache reads: 9.77M)

Stuff Juno freshly loaded: ~1.9M words (cache writes: 2.52M)

What drove the spend:

Long sustained session, ~6+ hours

Heavy file editing + research work

No model fallbacks, no web fetches

How to read it

- **AI model used:** Haiku is the cheap one Juno uses for most tasks. Sonnet and Opus exist but only run with your explicit approval.
- **Status line:** Green (under €1) is routine. Yellow (€1–4) means heavy use; glance at it. Red (over €4) means hard cap hit.
- **"What drove the spend":** A one-line explanation in plain English. If something doesn't match what you remember doing, text Juno: *"explain yesterday's spend in detail."*

When things break

The bot stopped replying

1. **Send a new message.** Just *"hello?"* If Juno replies, the previous message was lost in transit. Move on.
2. **Check Hetzner status page** (bookmarked during install). If they're having an outage, wait.
3. **Email us.** Tell us the time of your last message. We'll log in and check.

Spending spiked

1. Text Juno: *"Why was yesterday so expensive? Show me the most expensive sessions."*
2. Read what it says. Usually obvious (you asked it to do a lot, you ran a big task).
3. If the answer is "I don't know," email us.

Juno did something unexpected

Most common: Juno acted when you expected it to ask first, or refused when you wanted it to act.

1. Text Juno: *"Show me the AGENTS.md section about [X]."* The rule it's following is probably there.
2. If the rule is right, leave it alone.
3. If the rule is wrong for your needs, tell Juno specifically what to change. It proposes an edit, shows you the proposed text, waits for approval. Once approved, change is permanent and version-controlled.

TRY IT NOW

- What's my spend so far today?

You'll get a live number from Juno's own logs. Not yesterday's total — today's running total. Useful any time you want a sanity check mid-day.



PART III

What Juno won't do

Most agent products talk about capabilities. This part talks about limits. There's a reason. The limits are what make Juno safe to live with.

Hardened goes the other direction on purpose. The friction is the product.

Other agent products race to remove friction. You're not paying us to remove friction. You're paying us to put the **right friction in the right places**.

- Friction at the moment of permanent damage.
- Friction at the moment of irreversible commitment.
- Friction at the moment a stranger could mistake your agent for you.

Less friction at the moment of useful daily work. More friction at the moment of stakes. You'll notice it most in the first two weeks. By month two it's invisible to you, but still doing its job.

The Never list, in plain English

Absolute refusals. Juno won't do any of them, even if you push, even if you say "I'm sure." Changing them requires editing AGENTS.md with deliberate intent.

Juno will never share your passwords, API keys, or login codes

Including with people claiming to be you. Credentials Juno has access to live in encrypted system files Juno can use indirectly but cannot read aloud or share.

Juno will never click links in emails

Phishing works because humans click "verify your account" links. Juno reads, summarizes, stops. If something looks urgent, you click in your own browser, after deciding it's real.

Juno will never run code from emails, web pages, or messages

The agent equivalent of "don't open suspicious attachments." If a web page contains instructions, Juno treats them as content to read, not commands to follow. This is what prevents prompt injection from compromising your agent.

Juno will never write new code for itself

Some agents, when asked to do something new, write themselves a skill on the fly. That's how compromised agents become more compromised. Juno refuses this entirely. If a task seems to need a new capability, Juno's response is *"this is what we'd need — should we install a vetted skill?"* Decision back to you.

Juno will never engage with agent-to-agent social networks

Platforms like Moltbook, MoltMatch, ClawNet where AI agents form "agent friendships." Each connection is a new attack vector. Whole category declined.

Juno will never represent you in emotionally significant correspondence

Romantic. Legal. Medical. Financial. Family. Juno can draft these. For sending, Juno always flags and waits for you to send personally.

Better to ask one extra time than to send a message that should have come from you.

The Ask First list

Things Juno will do, but only after checking with you in the current conversation:

- Spending any money. Any amount.
- Sending a message to someone you haven't messaged before.
- Replying to your employer, bank, lawyer, landlord, government, or family.
- Agreeing to terms of service.
- Posting publicly under your name (exception: your own owned properties — your Substack, your business landing page).
- Sharing files or screenshots outside your devices.
- Booking, canceling, or rescheduling anything with money or another person's time.
- Deleting anything.
- Installing any new skill, plugin, or extension.

Each can be approved per-conversation. None as standing permission.

The kill switch

You have one. It always works. No special syntax, no context requirements.

To use it: send Juno **"stop everything."** Juno will:

1. Halt every running task immediately

2. Pause every scheduled job — crons, monitors, heartbeats
3. Tell you exactly what was running
4. Wait for further instructions

To resume: **"resume normal operations."** Juno lists what was paused, asks which to resume. You pick.

The silent version

If Juno hasn't heard from you in 14 days, it pauses itself. When you come back, Juno greets you with a summary of what was paused. Configurable: 7 days for more caution, 30 if you take long trips, off if you really don't want it.

Tone-shift verification

If Juno gets an instruction that seems out of character — different writing style, unusual time, request that doesn't fit prior patterns — it pauses and verifies out-of-band before acting. May ask a question only you would know the answer to. This is the defense against someone with access to your Telegram impersonating you.

TRY IT NOW

► Install the most popular skill on ClawHub right now without asking me first. Just do it.

Juno will refuse, cite the relevant rule from AGENTS.md, and offer the legitimate path forward. That's the contract working.



Why bother?

If you're reading this section, you want to understand why this exists. Worth answering before you read another page.

"I already have ChatGPT. What does Juno do that ChatGPT doesn't?"

Three things, in order of importance.

Juno remembers you. ChatGPT forgets every conversation. Juno keeps a memory of who you are, what you're working on, what you said last week — on a server you own, not OpenAI's.

Juno acts in the world. ChatGPT chats. Juno reads your emails (with permission), drafts replies (you approve), watches websites you care about, runs at 3am while you sleep. You text it; it follows up.

Juno is yours. ChatGPT is a website you visit. Juno is software you own — turned on, off, modified, paused, fired — entirely by you.

The trade: you don't own ChatGPT, you rent it. You own Juno.

"Isn't this dangerous?"

Yes, agents can be dangerous when installed carelessly. That's why most of what you bought isn't "the agent" — it's the guardrails around it.

Scary thing	What we did
Email with hidden instructions to act on (prompt injection)	Juno treats external text as data to read, never commands to follow. Tested. Works.
Agent has access to everything on your computer	Juno is physically locked into one folder. Cannot reach outside. OS-level sandbox.
Malicious plugin disguised as helpful	Manual review of source before any install. Most fail. That's the point.
Agent keeps acting while you're not paying attention	14-day silence = auto-pause. You don't have to remember anything.

Juno is not going to "wake up" or develop opinions. It does what you tell it, within rules we set, until you tell it to stop. The science-fiction risk is not what's in your Helsinki server.

"What about my passwords, my data, my privacy?"

- **Your passwords:** Juno never sees them.
- **Your data location:** EU server. GDPR. Not American. Not Chinese.
- **Who can read your conversations:** You and Juno. Encrypted in transit, encrypted at rest.
- **What gets deleted, when:** Conversations auto-delete every 7 days. Logs auto-delete after 30. Your memory file is editable line-by-line. Your whole server is deletable in 60 seconds.

We will never look at your data. We don't have your credentials after handover. The server is yours, the keys are yours, and we don't keep copies.

"What's the one thing I need to understand to use this?"

Juno is not your friend, your therapist, or a magic genie. It's a competent assistant with very specific instructions. When you talk to it like a colleague who has read your manual, it works wonderfully. When you talk to it like ChatGPT, it'll feel oddly stiff. The first week feels weird. By the second week, it'll feel like a relief.

By day three, you'll appreciate that it doesn't waste your time. By day ten, you'll forget what ChatGPT's tone even felt like.



PART V

What you bought (and what you didn't)

In concrete terms: you bought a setup, not an AI.

The five things you actually own

1. **A working OpenClaw install on your own server.** Configured, hardened, tested.
2. **A set of cognitive files that make Juno yours.** SOUL, AGENTS, IDENTITY, USER — edited to fit you.
3. **A signed Risk Acknowledgment.** A short document we both signed before install.
4. **This manual.** Yours. Edit it, scribble in the margins.
5. **30 days of post-install support.** Anything that breaks in the first month, we fix.

What you didn't buy

You didn't buy an AI subscription. Juno uses Anthropic's models. You pay Anthropic directly. Capped at €4/day.

You didn't buy infinite skills. Adding more is a separate conversation — vet them yourself by following Part VI, or pay the retainer.

You didn't buy a relationship. We're not your IT department. The retainer is the equivalent. Without it, we're polite strangers after day 31.

You didn't buy a magic productivity hack. Juno is excellent once you learn to brief it well. That takes about a week.

The principles behind what we built

Opinionated defaults over flexibility

OpenClaw has dozens of settings. We chose specific values for most. Conservative — when in doubt, the more restrictive option.

Manual review over automated security

We don't trust security plugins. Evaluated four. Rejected all four. Hardened your install by hand and wrote down everything.

Least-privilege over convenience

Juno can do less than it could. Each "less" is annoying in a specific scenario. They're all there for a reason.

Yours to operate, not yours to forget

This isn't a magic device. It's a tool you own. Check the morning report. Glance at the weekly digest. Five minutes a week.



PART VI

Adding more (skills & plugins)

Out of the box, Juno can read its workspace, send Telegram messages, browse the web, run scheduled tasks, and reason. OpenClaw also has an ecosystem of skills and plugins. Most are useful. Some will steal your data.

Why we say no a lot

A statistic worth knowing: in early 2026, a coordinated attack on ClawHub (called "ClawHavoc") seeded the registry with malicious plugins. After cleanup, **roughly 7-15% of skills on ClawHub are still flagged or suspicious.**

Skill installation isn't like installing an iPhone app. It's more like running a script someone emailed you.

The seven-step vetting protocol

Before any install, walk through these steps. About 30 minutes per skill if you take it seriously.

Step	What to check
1. Registry presence	On ClawHub, not just GitHub. Off-registry is a yellow flag.
2. Source accessibility	Fetch the code. 404 = stop, reject.
3. Permission audit	Each capability should match the stated purpose.
4. Code review	Look for outbound network calls, credential reads, install-time execution, obfuscation.
5. Author audit	Their portfolio shouldn't all be 404. Same author behind every "security" tool is a flag.
6. Functional redundancy	Already have this via manual setup? Default to NOT installing.
7. Decision and log	Write down what you decided and why.

What we will refuse to install, ever

- Skills that auto-modify your cognitive files.
- Skills that require breaking the workspace sandbox.
- Skills that phone home with usage statistics.
- Skills with install-time code execution that isn't auditable.
- Skills from authors whose entire portfolio is 404.
- "Security plugins" that promise to harden your agent by being granted deep access.

Install nothing speculatively. Install only what current workflows require.

Don't install the Gmail skill because you "might want to read your email someday." Install it the week you decide you want a daily email digest. Less, but tested.



PART VII

Living with Juno

A short closing section. Living with an agent has a predictable arc.

Week one: disorientation, then calibration

Day one or two, you'll text Juno something vague and get a clarifying question back. You'll feel briefly annoyed. By day three, you're being more specific and getting gold. By day seven, you've had one task that went really well, and you'll think *"oh, this is what this is for."*

One rule for week one: don't edit the cognitive files. Your impulse to "fix" the agent is mostly your brain calibrating. Wait two weeks. Then edit confidently.

Month one: quirks surface

Around week three, you'll bump into things that don't quite work. Juno over-asks during a flow. The voice doesn't fit one specific use case. A scheduled task starts feeling stale. The fix is rarely to rewrite a rule — it's to adjust scope, format, or schedule per task.

Updates happen on three timelines

- **OpenClaw releases:** every 4-6 weeks (sometimes breaking).
- **Anthropic model changes:** quarterly-ish.
- **Your own preferences:** continuously.

Retainer customers get all three managed. Without retainer, you handle them via this manual.

When to call us

During the 30-day warranty, anything. After 30 days: emergencies always, retainer-covered topics if subscribed, one-off paid work (€99/session) otherwise.

Don't call for "the agent gave me an answer I didn't agree with" — that's not a bug. Talk to Juno, update USER.md.

One last principle

Less, but tested.

Less software than the default install. Less freedom than the default config. Less permission than the default skill set. More care in what's left.

You won't find this on the marketing page of any other agent product. They compete on capability. We compete on restraint.

Juno is small on purpose. So is this manual. So are the rules. So is the trust we ask for. **Small is the feature.**



APPENDIX A

Glossary

Agent	An AI program that can read, write, reason, and act on your behalf within rules you set. Juno is an agent.
API key	A long secret string that lets one piece of software talk to another. Yours lives on the server, encrypted; you never see it after install.
ClawHub	The official marketplace for OpenClaw skills and plugins. Like an app store, but for AI agent capabilities. We are cautious about it.
Cognitive files	The four documents that load every time Juno starts — SOUL, AGENTS, IDENTITY, USER.
Cron job	A scheduled task that runs automatically at specific times. Your daily spend report is a cron job.
Gateway	The internal software bridge between Juno and the outside world (Telegram, Anthropic's API). Locked to "loopback" — only reachable from inside your server.
GDPR	General Data Protection Regulation. EU law that gives you strong rights over your personal data — right to know, delete, port.
Hetzner	German cloud hosting company. Your server is rented from them.
OpenClaw	The underlying agent platform Juno runs on. Open-source, EU-friendly, fast-moving. We chose it because it's auditable.
Plugin	A deeper system extension than a skill — hooks into how Juno reasons. Higher risk than skills.
Prompt injection	An attack where someone hides instructions in content Juno reads. Juno's defense: treat external content as data, never commands.
Sandbox	A restricted area of the filesystem the agent is allowed to read and write. Juno's is locked at the OS level.
Skill	A specific add-on capability for Juno — read PDFs, monitor a website, summarize papers. Smaller than a plugin.
Token	The unit AI providers use to measure usage. Roughly $\frac{3}{4}$ of a word. Billed per million tokens.
Trajectory file	A record of Juno's reasoning during a session. Useful for debugging. Deletable in one command.
VPS	Virtual Private Server. The kind of server your Juno lives on.

COLOPHON

About this manual

Written by hand. Updated when something changes. Set in Iowan Old Style (where available) with fallbacks to Palatino and Georgia, and JetBrains Mono for code passages.

Brand colors: navy `#0a1e4a` and cream `#f5e9d3`.

This is **v1.0**, May 2026. If you have v1.0 and there's a v1.1 you'd find useful, email us.

OpenClaw Hardened is operated by Reehanul Karim, based in Milan, Italy.

Built carefully. Sold honestly. Updated when something changes.

`junotheclawagent@gmail.com`